

Random graph generation

The generation of connected random graphs is non-trivial and important to many applications. In particular, it is not easy to sample a connected random graph from the space \mathcal{G} of all connected graphs with uniformity; i.e. without a bias for graphs of a special kind. This entry contains three algorithms for sampling random connected graphs. The first two are simple and of reasonable complexity, but biased. The third algorithm is less efficient, but it can be proven that it samples graphs without bias.

Bottom-up approach

The bottom-up approach consists in generating a simple graph and constructing a random graph from it. We shall use spanning trees to (surprise) span graphs from them. Some definitions are in order.

- Let \mathcal{T}_n the set of all trees of n vertices, \mathcal{G}_n the set of all graphs with n vertices, and $\mathcal{G}_{n,m}$ the set of all graphs with n vertices and m edges. We shall assume the vertices of these graphs are labeled $1, \dots, n$.
- For any $T \in \mathcal{T}_n$, we define $\mathcal{U}_T := \{G \in \mathcal{G}_n : T \subseteq G\}$ and refer to it as *the universe* of T . Thus, the universe of $T \in \mathcal{T}_n$ is the set of graphs spanned by T .
- Let $\Lambda(n) = \{x, y : x, y \in 1, \dots, n\}$, the set of all possible edges in a graph of n vertices.

Before discussing any graph generation, I must review an important concept relating to trees: the Prüfer sequence.

Prüfer sequence

The Prüfer sequence of $T \in \mathcal{T}_n$ is a word over the alphabet $\Sigma = \{0, \dots, n-1\}$ of length $n-2$. Prüfer proved that there is a bijection between \mathcal{T}_n and Σ^{n-2} (which incidentally provided a nice proof of the fact that there are n^{n-2} distinct trees of n vertices).

The algorithms hereby presented generate random trees by generating random Prüfer sequences. It is easy to construct algorithmically the tree corresponding to a Prüfer sequence, and thus a random tree is obtained.

Let T_2 denote the unique tree of two vertices. Let L be the label-set, i.e. the set of natural numbers which label the vertices of our graphs. Then, a recursive algorithm for constructing the corresponding to a Prüfer sequence $p = p_1 \dots p_{n-2}$ over a label set L is the following:

begin align textbf begin textbf f(p, L) quad quad textbf f|p| = 0 textbf do quad quad quad quad textbf return T_2 textbf*

Here is an illustration of the algorithm ran on $p = 17375$, $L = \mathbb{N}_7$. The illustration is taken from this source.

Random graph generation

Let T_w denote the graph corresponding to the Prüfer sequence w . Given a tree T we define a special family of edges:

$$S_T = \{e \in \Lambda(n) : e \notin E(T)\} = \Lambda(n) - E(T)$$

These are called the *spanning edges*, since these are the edges required to span connected graphs from the spanning tree T . Observe that $\emptyset \in S_T$ and is the set of edges required to span T out of T .

Then, for any fixed n , we let the language $1, \dots, n^{n-2}$ be the index set of an indexed family of functions \mathcal{F} defined as:

$$\mathcal{F}(w) : \mathcal{U}_{T_w} \rightarrow S_T$$

$$\mathcal{F}(w)(G) = E(G) - E(T_w)$$

$\mathcal{F}(w)$ is a bijection. Let $S \in S_{T_w}$ for an arbitrary T_w . Define $G = (V(T), S \cup E(T_w))$. Then $T_w \subseteq G$ and G belongs to the domain of $\mathcal{F}(w)$. $E(T_w) \cap S_{T_w} = \emptyset$ by definition. Then $S = E(G) - E(T_w)$.

$\therefore \mathcal{F}(w)$ is surjective.

Let $G, G' \in \mathcal{U}_{T_w}$ for an arbitrary T_w . Assume $\mathcal{F}(w)(G) = \mathcal{F}(w)(G')$. Assume $G \neq G'$. Since T_w spans both, all edges in T_w are in G, G' and their difference must lie in an edge outside of T_w . But all edges outside of T_w are in $\mathcal{F}(w)(G)$ and $\mathcal{F}(w)(G')$ respectively, which are the same. This is a contradiction. $\therefore G = G'$.

$\therefore \mathcal{F}(w)$ is injective.

$\therefore \mathcal{F}(w)$ is bijective.

To summarize the relations we have established, we see that there is a nice generative path

$$\text{Prüfer sequence} \rightarrow \text{Tree} \rightarrow S \in S_T \rightarrow G \in \mathcal{U}_T$$

which inspires effective procedures for random graph generation. The most obvious procedure is the following. Given an desired number of vertices n ,

- (1) Generate randomly $p = p_1 \dots p_{n-2} \in \Sigma^{n-2}$.
- (2) Span the tree $T = (V, E)$ of the Prüfer sequence p .
- (3) Let $k \in_R 0, \dots, \frac{n(n-1)}{2}$.
- (4) Let $\ell_1, \dots, \ell_k \in_R S_T$, all distinct.
- (5) Let $E = E \cup \ell_1, \dots, \ell_k$

Because all trees of n vertices correspond to a sequence, all trees can be sampled. And all connected graphs can be derived from the set of all spanning trees. Then this procedure generates all graphs in \mathcal{G}_n .

The question is whether it is equally likely to generate any two graphs in \mathcal{G}_n . It is obvious that it is equally likely to generate any tree. And the probability that a given graph is generated depends entirely on the number of spanning trees it contains. Not all graphs have the same number of spanning trees. \therefore It is more likely to generate a graph with many spanning trees than a graph with few spanning trees.

The second effective procedure extends the input from only n , the number of vertices, to m , the number of edges. Thus, it specifies the problem further into the question of how to generate more or less dense graphs of n vertices. The effective procedure consists simply in generating T_w and sampling $\ell \in_R S_T$ repeatedly until $|E_T| = m$.

textbfInput : n, m

$T := (V, E) =$
textbfgenRandomTree(n)

$S_T := [$
 $\textit{Gamma}^c(v_1),$
 $\textit{ldots},$
 $\textit{Gamma}^c(v_n)]$

$C := [|S_T[1]|,$
 $\textit{ldots}, |S_T|[n]]$

$V := [1,$
 $\textit{ldots}, n]$

$n' := n$

textbfwhile $|E(T)| < m$
textbfdo

qquad $:=$
textbfrandom($1, n'$)

qquadv $:= V[i]$

qquad
textbfif $d(v) == n - 1$
textbfdo

qquadqquadtextbfdeleteElement

qquad
qquad
textbfdeleteAt(V, i)

qquad
qquadn' $:= n' - 1$

Generating a tree from a random Prüfer sequence is $O(n^2)$. Forming S_T is also $O(n^2)$. Within the while loop there are simply index manipulations, so the complexity of the loop is $\varphi \times O(1) = \varphi$, with φ the complexity of the number of iterations.

All iterations add an edge except those where a saturated vertex is chosen. A saturated vertex may be chosen at most once. \therefore There are $O(n)$ iterations where a saturated vertex is chosen. Since the rest of the iterations add an edge, their number is fixed: $m - (n - 1)$, where $(n - 1)$ is the number of edges in the spanned tree. \therefore There are exactly $m - n + 1$ edge-adding iterations.

$$\therefore \varphi = O(n) + O(m - n + 1) = O(n) + O(m) = O(m)$$

$$\therefore \text{The complexity of the algorithm is } O(n^2) + O(m) = O(n^2).$$

I showcase here random graphs with their source random trees as generated by this algorithm. The left-most graph is the random tree which spanned the right-most graph. The algorithm was implemented in C but the generated graphs were plotted using the `networkx` Python package.

Top-down approach

A top-down approach analogous to the previous algorithms would in principle consist in the generation of a K_n and the random removal of edges in the graph. This involves some extra complexity: while adding edges to a connected graph cannot disconnect it, removing edges from a graph may do so. Thus, on each iteration, some algorithmic process should determine if an edge can be deleted or not without violating the connectivity invariant.

Thus, an effective procedure goes as follows. Given n and a desired number of edges m ,

- (1) Generate $K_n = (V, E)$.
- (2) Let $V_c = V$ the list of vertices which can be pruned.
- (3) Let $v \in_R V_c, w \in_R \Gamma(v)$.
- (4) Check if removing v, w preserves connectivity; if so, remove it from E .
- (5) Define $\Gamma(v) = \Gamma(v) - w, \Gamma(w) = \Gamma(w) - v$, regardless of whether the edge was removed.
- (6) If $d(z) = 1$ or $\Gamma(z) = \emptyset$, remove z from V_c , where $z \in v, w$.
- (7) If $|E| \neq m$, go to (3).

Assume **BFSFind**(E, x, y) is an algorithm that checks if there is a path from x, y in a given edge set. Then a pseudo-code implementation of the effective procedure above is:

$$\begin{array}{l} (V,E):=\\ \textit{textbf{genKn}}(n) \end{array}$$

$$\begin{array}{l} \textit{Gamma}:= [\\ \textit{Gamma}(v_1),\\ \qquad \qquad \qquad \textit{ldots},\\ \textit{Gamma}(v_n)] \end{array}$$

$$\begin{array}{l} V_c := [v_1,\\ \qquad \qquad \qquad \textit{ldots},v_n] \end{array}$$

$$\begin{array}{l} \textit{textbf{while}}|E|\\ \qquad \qquad \qquad \textit{neqm}\\ \textit{textbf{do}} \end{array}$$

$$\begin{array}{l} \textit{quad}\\ \textit{quadv} =\\ \textit{textbf{frandFrom}}(V_c) \end{array}$$

$$\begin{array}{l} \textit{quad}\\ \textit{quadv} =\\ \textit{textbf{frandomFrom}}(\\ \qquad \qquad \qquad \textit{Gamma}[v]) \end{array}$$

$$\begin{array}{l} \textit{quad}\\ \textit{quad}\\ \textit{Gamma}[v] =\\ \textit{Gamma}[v]-\\ \qquad \qquad \qquad w \end{array}$$

$$\textit{quad}$$

$$\textit{quad}$$

$$\textit{Gamma}[w] =$$

$$\textit{Gamma}[w] -$$

$$v$$

$$\textit{quad}$$

$$\textit{quad}$$

Generating a K_n is $O(n^2)$. The **while** selects a random vertex from V_c and attempts to erase one of its edges. There is only one case one an edge is not removed; namely, when the sampled edge is a bridge. This happens at most once per bridge. There are at most $n - 1$ bridges in a graph. \therefore There are $O(n)$ iterations which do not remove an edge.

The remaining iterations will remove an edge and there will be exactly $\frac{n(n-1)}{2} - m$ of them.

\therefore There are $O(n) + O(\frac{n(n-1)}{2} - m) = O(n^2 - m)$ iterations.

The operations in each iteration are $O(1)$ except BFS. The complexity of BFS grows linearly with the number of edges. \therefore BFS is $O(n^2)$.

\therefore The algorithm is $O(n^2) + O(n^2 - m)O(n^2) = O(n^4 - n^2m)$.

In practice the algorithm will perform better than this. BFS stops whenever w is found starting from v . This still is asymptotically $O(|E|)$, but in practice the bound will seldom be reached. Furthermore, BFS is ran on increasingly sparser graphs. Its asymptotic complexity is given by the number of edges in the initial K_n , but it decreases with each edge-removing iteration.

Below, I display a K_{100} and the random tree generated from it.

To prove that our algorithm is correct and unbiased requires more formalization. Let $\mathcal{C}_{n,m} \subset \mathcal{G}_{n,m}$ be the set of connected graphs of n vertices, m edges. Let $\mathbb{C}(n,m) = |\mathcal{C}_{n,m}|$. Let $\mathcal{E}_{n,m}$ be the class of edges which, if removed from a K_n , produce a graph in $\mathcal{C}_{n,m}$. Since any $G \in \mathcal{C}_{n,m}$ corresponds univocally to a set of edges $W \in \mathcal{E}_{n,m}$,

$$|\mathcal{E}_{n,m}| = \mathbb{G}(n,m)$$

and there is a bijection $f_{n,m} : \mathcal{E}_{n,m} \mapsto \mathcal{C}_{n,m}$ mapping a set of edges to the connected graph generated by removing those edges from K_n .

We shall prove that (1) our algorithm effectively constructs a $W \in \mathcal{E}_{n,m}$ and computes $f(W)$ and (2) that any $W \in \mathcal{E}_{n,m}$ has an equal probability of being constructed.

(1) The algorithm iteratively removes edges ensuring that the connectivity invariant is preserved. It is trivial to see that it removes $k := \binom{n}{2} - m$ edges. Let $S = e_1, \dots, e_k$ be the set of randomly sampled edges, where e_i was sampled at the i th edge-removing iteration.

Preserving the invariant entails that, across iterations, the sampling spaces E_1, \dots, E_r from which edges to remove are sampled are:

- $E_1 =$
 $e \in W : W \in \mathcal{E}_{n,m}$
- $E_2 =$
 $e \in W : W \in \mathcal{E}_{n,m} \wedge e_1 \subseteq W$
- $E_3 =$
 $e \in W : W \in \mathcal{E}_{n,m} \wedge e_1, e_2 \subseteq W$

etc. Thus, the general form is $\mathcal{E}_i = \{e \in W : W \in \mathcal{E}_{n,m} \wedge \{e_1, \dots, e_{i-1}\} \subseteq W\}$.

It follows that $S =$

$e_1, \dots, e_k \subseteq W$ for some $W \in \mathcal{E}_{n,m}$. But $|S| = |W| = k$. Then $S = W$ and $S \in \mathcal{E}_{n,m}$. And since S is the set of removed edges, the algorithm computes $f(S)$.

(2) Since there is a bijection between $\mathcal{C}_{n,m}$ and $\mathcal{E}_{n,m}$, a graph is more probable than others if and only if there is a set $S \in \mathcal{E}_{n,m}$ that is more probably constructed than others. This could only be true for two cases: (1) An edge or set of edges in S is more likely to be chosen, or (2) S contains more elements than other members of $\mathcal{E}_{n,m}$. But (1) is impossible if the selection is random, and (2) contradicts that $|S| = \binom{n}{2} - m$ for every $S \in \mathcal{E}_{n,m}$.

\therefore The algorithm is correct and unbiased.

As an appendix, let us derive $\mathbb{C}(n, m)$. It is straightforward to reason that

$$|\mathcal{G}_{n,m}| = \binom{\frac{n(n-1)}{2}}{m}$$

This inspires the definition of a generating function for the class $\mathcal{G}_{n,m}$ of graphs of n vertices, m edges:

$$\begin{aligned} A(z, u) &= \sum_{k=0}^{\infty} \left(\sum_{r=0}^{\infty} \binom{\frac{k(k-1)}{2}}{r} u^r \right) \frac{z^k}{k!} \\ &= \sum_{k=0}^{\infty} (1+u)^{\frac{k(k-1)}{2}} \frac{z^k}{k!} \\ &= 1 + \sum_{k=1}^{\infty} (1+u)^{k(k-1)/2} \frac{z^k}{k!} \end{aligned}$$

An important observation is that any $G \in \mathcal{G}_{n,m}$ is a set of elements $G_1, \dots, G_r \in \mathcal{C}_{n,m}$. Informally, any graph is a set of connected graphs. Since the relationship of the class $\mathcal{G}_{n,m}$ and the class $\mathcal{C}_{n,m}$ is the set-of relation, the generating function $C(z, u)$ of $\mathcal{C}_{n,m}$ satisfies

$$A(x) = e^{C(z,u)}$$

Then

$$\begin{aligned}
C(z, u) &= \ln \left[1 + \sum_{k=1}^{\infty} (1+u)^{k(k-1)/2} \frac{z^k}{k!} \right] \\
&= \sum_{q=1}^{\infty} \frac{(-1)^{q+1}}{q} \left[\sum_{k=1}^{\infty} (1+u)^{k(k-1)/2} \frac{z^k}{k!} \right]^q
\end{aligned}$$

Thus, $C(z, u)$ is the generating function of the composite sequence $|C_n, m|_{n \geq 1, m \geq 0}$. The resulting expression for the number of connected graphs of n vertices, m edges is:

$$\mathbb{C}(n, m) = n! [z^n][u^m] \sum_{q=1}^{\infty} \frac{(-1)^{q+1}}{q} \left[\sum_{k=1}^{\infty} (1+u)^{k(k-1)/2} \frac{z^k}{k!} \right]^q$$

where $[z^n][u^m]C(z, u)$ are the n th and m th coefficients of the generating functions for the polynomials dependent on z and u , respectively.

Example. For $m = n-1$ (i.e. connected trees) across $n = 2, 3, \dots$, \mathbb{C} effectively produces the sequence

$$1, 1, 3, 16, 125, 1296, \dots = \{n^{n-2}\}_{n \geq 2}$$

which we know is correct because there are n^{n-2} Prufer sequences of length n .