

Random Graph Generation

Santiago López Pereyra

Introduction

The generation of connected random graphs is non-trivial and important to many applications. In general, given $n, m \in \mathbb{N}$, it is not difficult to sample a random graph from the space of all graphs of n vertices, m edges. The problem becomes more difficult when we require (a) that the randomly generated graph be *connected* and, if possible, (b) that any possible such graph has the same probability of being generated (i.e. that we sample the connected graphs with uniformity).




Figure 1: Randomly generated graphs

A direct and simple algorithm is to generate a random tree (e.g. via a Prufer sequence) and

add edges randomly until the desired number of edges is reached. This procedure is relatively efficient ($\mathcal{O}(n^2)$), but it is biased. Not all connected graphs have the same number of spanning trees, and therefore the probability of generating a given connected graph is not uniform.

A computationally heavier approach is to generate a K_n and prune edges randomly until the desired number of edges is reached, while maintaining the connectivity invariant. This algorithm is unbiased, but its complexity is higher. The procedure is as follows:

```

(V, E) := genKn(n)
Ec := [e1, ..., e|E|]
while |E| > m do
    {v, w} := randSample(Ec)
    E := E - {v, w}

    if ¬ConnectivityCheck(E, v, w) then
        E := E ∪ {v, w}
        Ec := Ec - {v, w} // edge is a bridge
    fi
od
return (V, E)

```

Generating a K_n is $\mathcal{O}(n^2)$. The **while** selects a random edge from E_c and attempts to prune it. There is only one case in which an edge is not removed; namely, when the sampled edge is a bridge. This happens at most once per bridge. There are at most $n - 1$ bridges in a graph. Hence, there are $\mathcal{O}(n)$ iterations which do not remove an edge.

The remaining iterations will remove an edge and there will be exactly $\frac{n(n-1)}{2} - m$ of them.

∴ There are $\mathcal{O}(n) + \mathcal{O}(\frac{n(n-1)}{2} - m) = \mathcal{O}(n^2 - m)$ iterations.

The operations in each iteration are $\mathcal{O}(1)$ except the connectivity check, which we assume to be a BFS search with starting vertex v and target vertex w . BFS is $\mathcal{O}(n^2)$.

∴ The algorithm is $\mathcal{O}(n^2) + \mathcal{O}(n^2 - m)\mathcal{O}(n^2) = \mathcal{O}(n^4 - n^2m)$.

In practice the algorithm will perform better than this. BFS stops whenever w is found starting from v . This still is asymptotically $\mathcal{O}(|E|)$, but in practice the bound will seldom be reached. Furthermore, BFS is ran on increasingly sparser graphs. Its asymptotic complexity is given by the number of edges in the initial K_n , but it decreases with each pruning iteration.

To prove that the algorithm is unbiased we need a few definitions. Let $\mathcal{G}_{n,m}$ denote the set of all graphs with n vertices and m edges, and let $\mathcal{C}_{n,m} \subset \mathcal{G}_{n,m}$ denote the subset of *connected* graphs.

Let $\mathcal{E}_{n,m}$ be the class of edge sets $W \subseteq E(K_n)$ such that removing W from K_n produces a connected graph with m edges. Each $G \in \mathcal{C}_{n,m}$ corresponds uniquely to one such $W \in \mathcal{E}_{n,m}$, so

$$|\mathcal{E}_{n,m}| = |\mathcal{C}_{n,m}|.$$

It follows that there is a bijection

$$f_{n,m} : \mathcal{E}_{n,m} \rightarrow \mathcal{C}_{n,m}, \quad f_{n,m}(W) = (V, E(K_n) - W).$$

We now prove that:

- (1) The edges removed by the algorithm form a valid set $W \in \mathcal{E}_{n,m}$ and the resulting graph is $f_{n,m}(W)$.
- (2) Each $W \in \mathcal{E}_{n,m}$ may be formed with equal probability.

(1) The algorithm removes $k = \binom{n}{2} - m$ edges $S = \{e_1, \dots, e_k\}$. The connectivity invariant is preserved at each successful removal, so $\{e_1\}, \{e_1, e_2\}, \dots, \{e_1, \dots, e_k\}$ are all members of $\mathcal{E}_{n,m}$. By construction, the final graph has edges $E(K_n) - S$, i.e. the final graph equals $f_{n,m}(S)$. ■

(2) Unbiasedness (symmetry argument) The algorithm removes the edges of some $W \in \mathcal{E}_{n,m}$ and each successful run corresponds to an ordered sequence (e_1, \dots, e_k) which is a permutation of the elements of that W . Note that:

1. If $W \in \mathcal{E}_{n,m}$, then for any subset $S' \subseteq W$ the graph $f(S')$ is still connected (removing fewer edges cannot disconnect a graph that remains connected after removing the larger set). Hence any permutation of the k edges of W is an admissible removal order: none of those edges will ever be a bridge at the moment it is removed. Therefore the number of admissible orders that produce W is exactly $k!$.
2. By vertex- and edge-symmetry of K_n , the size $|E_i|$ of the candidate set at step i depends only on i (and on n, m), not on which particular edges form S_{i-1} . Denote $a_i := |E_i|$.

At step i the algorithm chooses uniformly from E_i , so the probability of that exact ordered sequence is $\prod_{i=1}^k \frac{1}{a_i}$. Since there are $k!$ ordered sequences (permutations) that yield the same unordered set W , the probability of producing W equals

$$\Pr[\text{output } W] = k! \prod_{i=1}^k \frac{1}{a_i}.$$

The right-hand side does not depend on W (only on k and the sequence a_i), so every $W \in \mathcal{E}_{n,m}$ is produced with the same probability. Thus the induced distribution on $\mathcal{C}_{n,m}$ is uniform, i.e. the algorithm is unbiased.




Figure 2: A complete graph K_{100} (left) and a random tree obtained by pruning (right).

1 As a Markov chain

Let G_i denote the graph after i successful iterations of the edge-pruning algorithm. Then the sequence $\{G_i\}_{i \geq 0}$ defines a discrete-time Markov chain (DTMC) on the state space $\mathcal{C}_{n,m}$, the set of connected graphs with n vertices and m edges.

Markov property

The Markov property holds because the choice of edge to remove at step $i + 1$ depends solely on the current graph G_i and the remaining candidate edges E_c . Formally, for all $i \geq 0$ and all $G_0, \dots, G_{i+1} \in \mathcal{C}_{n,m}$,

$$\Pr[G_{i+1} \mid G_i, G_{i-1}, \dots, G_0] = \Pr[G_{i+1} \mid G_i].$$

Transition probabilities

Let $P(G_i \rightarrow G_{i+1})$ denote the one-step transition probability from graph G_i to graph G_{i+1} , assuming $G_{i+1} \neq G_i$. Then

$$P(G_i \rightarrow G_{i+1}) = \begin{cases} \frac{1}{|E_i|}, & \text{if } G_{i+1} \text{ can be obtained by removing a single non-bridge edge from } G_i, \\ 0, & \text{otherwise,} \end{cases}$$

If a sampled edge is a bridge, then the graph remains unchanged and $G_i = G_{i+1}$. The probability of this self-loop transition is given by

$$P(G_i \rightarrow G_i) = \frac{|B_i|}{m_i},$$

where B_i is the set of bridge edges at the current iteration and m_i is number of edges in G_i . Note that $B_i = \overline{E_i}$, which means $|B_i| = m_i - |E_i|$, entailing that

$$P(G \rightarrow G) = \frac{m_i - |E_i|}{m_i} = 1 - \frac{|E_i|}{m_i}$$

From this follows that the probability that an edge is successfully removed at iteration i is $|E_i|/m_i$.

Note that there are at most $n - 1$ bridges in a connected graph with n vertices, and this bound remains constant across iterations, so at any given point in time we have

$$P(\text{A bridge is chosen at iteration } i) \leq \frac{n - 1}{m_i}$$

Obviously, this bound is informative only when $m_i \gg n - 1$, since otherwise the bound approximates 1 and becomes trivial. This of course means that the bound is informative for dense graphs and not for sparse ones.

Stationary distribution

Since the algorithm is unbiased, every connected graph $G \in \mathcal{C}_{n,m}$ is equally likely to appear as the final output. This implies that the stationary distribution π of the Markov chain is uniform over $\mathcal{C}_{n,m}$:

$$\pi(G) = \frac{1}{|\mathcal{C}_{n,m}|}, \quad \forall G \in \mathcal{C}_{n,m}.$$

Irreducibility and aperiodicity

The chain is irreducible in the sense that, starting from K_n , any connected graph with m edges can eventually be reached by a sequence of valid edge removals. The chain is aperiodic because there is a positive probability of staying in the same state (when a bridge is sampled), i.e.,

$$P(G \rightarrow G) > 0 \quad \forall G \in \mathcal{C}_{n,m}.$$

Hence, the Markov chain converges to the uniform stationary distribution.

2 An improvement

The probability that a chosen edge is a bridge increases as the graph becomes sparser, but is relatively negligible in the early stages of the algorithm. Since such probability is bounded at each iteration by $n - 1/m_i$, and both of these quantities are known to the algorithm, we can define a tolerance ϵ s.t. if $n - 1/m_i < \epsilon$, we skip the connectivity check and automatically remove the sampled edge.